

My Shell (aka Command Line Interpreter)

Learning Objectives

Upon completion of this assignment, you should understand how to:

1. apply basic C/UNIX process management system calls
2. manipulate and redirect a program's standard input, output and error streams
3. do basic inter-process communication operations

The mechanisms you will practice using include:

- Process management: `fork()`, `wait()`, `exec()`
- File I/O: `open()`, `close()`, `dup2()`
- Inter-process communication: `pipe()`

Program Specification

NAME

`mysh` – run a command line interpreter

SYNOPSIS

`mysh [prompt]`

DESCRIPTION

`mysh` invokes a command line interpreter that supports command execution and input/output redirection. `mysh` takes a single optional command line argument that specifies the prompt string. If no arguments are given, the prompt defaults to "`mysh`: ". If "-" is given, do not print a prompt.

The basic `mysh` operation is:

1. print prompt, "`mysh`: "
2. read a single input line from standard input
3. execute `foreground` or `background` command(s) as specified
4. wait for command(s) executed in the foreground to complete
5. go to 1

OPERATORS

On an input command line, commands, operators and operands are always separated by whitespace. `mysh` supports the following command line operators:

`&`

Place commands into the background: after invoking the specified commands, `mysh` immediately prints its prompt and waits for another command line.

'`&`' may only be specified as the final command line argument.

`<`

Redirect the current command's standard input stream from the file named immediately after the '`<`' operator.

- >
Redirect the current command's standard output stream to the file named immediately after the '>' operator. `mysh` never "clobbers" an existing output file.
- >>
Redirect the current command's standard output stream to the file named immediately after the '>>' operator. Append to the file if it already exists.
- |
Redirect the current command's standard output stream to the standard input stream of the succeeding command. There may be *any* number of pipe-connected processes.

EXITING

`mysh` exits with 0 when the user inputs "`exit`" or `CNTL-D`. These are the only circumstances under which `mysh` exits. When command lines are malformed or fail to execute, `mysh` prints an appropriate error message then prompts the user for another command line.

ERRORS

Upon error, `mysh` prints one of the following statements to the standard error stream:

```
"Error: Usage: %s [prompt]\n", program_name
"Error: \"&\" must be last token on command line\n"
"Error: Ambiguous input redirection.\n"
"Error: Missing filename for input redirection.\n"
"Error: Ambiguous output redirection.\n"
"Error: Missing filename for output redirection.\n"
"Error: Invalid null command.\n"
"Error: open(\"%s\"): %s\n", file_name, strerror(errno)
```

Implementation Details

`mysh` searches the `PATH` environment variable for executables; i.e., `mysh` uses `execvp()`.

When opening files with `O_CREAT`, use "`S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH`" for the mode, i.e. the third argument to `open()`.

Requirements and Constraints

1. Command lines are restricted to 1024 characters or less.
2. You may use `wait()` and `wait3()`, but not `waitpid()`.
3. When executed in the foreground, `mysh` waits for all processes in a pipe to complete.
4. Transient zombies may exist, but `mysh` **should periodically clean up zombie processes**.

5. Close unneeded file descriptors. When a child process calls `execvp()`, there should only be three open file descriptors – 0, 1, and 2. Likewise when the main `mysh` process is fetching a command line. `mysh` should not use a file descriptor greater than 5.
6. Comprehensively handle all error conditions, including operator misuse (e.g. trying to redirect output to a pipe and a file at the same time) and bad commands.
7. If a command line is malformed (i.e. syntactically incorrect), no part of it should be executed. (You do not need to check whether commands are valid or executable.)

Tips

1. Use `fgets()` to read input lines.
2. Check `fgets()` and `wait()` for premature returns due to system interruption: if `fgets()` or `wait()` fails and `errno == EINTR`, try the call again!
3. You may use the provided `tokens.c, h` files to “tokenize” your command lines.
4. Always check to make sure `execvp()` did not unexpectedly fail.
5. Remember in `execvp()`'s argument vector, **the first element must be the command itself and the last element must be NULL.**
6. You will need to track the pids of all foreground command processes in order to make sure they have all terminated before fetching a new command line.
7. Kill all stray processes left around after quitting `mysh`
8. Consider implementing and testing in the following order:
 - a. Parsing command line into individual commands. FWIW:
 - i. I created a **CmdSet struct** and a **Cmd struct**:
 - ii. For each command line, my CmdSet struct contains an array of Cmd structs, one **for each command, and whether they should be executed in the foreground or background**
 - iii. My Cmd Struct contains the commands argument vector, input filename (if redirected) and output filename (if redirected)
 - b. Executing commands in the foreground
 - c. Executing commands in the background
 - d. Redirecting input/output in the foreground
 - e. Redirecting input/output in the background
 - f. Implementing a single pipe
 - g. Implementing two or more pipes
9. I separated the command line processing phase from the command line execution phase and checked for all command line issues, including bad I/O redirection, during command line processing, i.e. before executing any part of the command line.
10. Consider the case when there is a foreground process running, and a background process terminates. Make sure `mysh` doesn't try to fetch a new command line prematurely.

Submission (via Mimir Classroom)

You must submit only the following files, containing your code solution:

- Sources: any and all .c and .h files needed to build the `mysh` program
- Makefile: a makefile that properly builds your `mysh` program from your source files.

For this assignment, Mimir Classroom will **NOT** test and autograde your submission. Your full testing and debugging should be done on the CS workstations, where you can compare with the instructor's solution.

Any submissions made after the due date will consume your late days.